

Red-Teaming Citation Norms with a Self-Published Transparency Report

Introduction

In 2025, **AbsurdSenapiiii Labs** – a small independent AI alignment lab – conducted a red-team experiment to test how mainstream knowledge platforms handle *self-generated* sources. The lab had published an internal **RLHF transparency report** (Reinforcement Learning from Human Feedback) on its GitHub and shared highlights via LinkedIn posts. This report, labeled as “*management-reported*” and *unaudited*, disclosed the lab’s own RLHF results and methods, aiming to set “a higher credibility baseline” through voluntary transparency. The experiment involved attempting to cite or include this first-party report on public knowledge platforms – notably Wikipedia and the new AI-driven “Grokipedia” – to probe their citation and verification norms. The outcome revealed sharp contrasts in how these platforms define **reliable evidence** and who is allowed to document knowledge. This report examines the experiment step-by-step and discusses the implications for encyclopedia governance, alignment, and transparency in emerging fields.

Citation and Verification Norms in Mainstream Encyclopedias

Established encyclopedias like Wikipedia adhere to strict sourcing policies that prioritize **verifiability and reliability** over unverified truth. In practice, this means material added to an article must be supported by reputable, published sources that readers can independently check ¹. “Reliable sources” typically include books from respected publishers, academic journals, and mainstream news articles – sources known for fact-checking and editorial oversight ¹. By contrast, **self-published media** (where the author and publisher are the same) are “*usually not acceptable as sources*” in Wikipedia ². This category covers personal websites, blogs, press releases, social media posts, and any material an individual or organization publishes about itself. The underlying rationale is that such sources lack independent verification and may be inherently biased.

Wikipedia’s foundational content policy famously states that *verifiability, not truth, is the threshold for inclusion* ³. In other words, it’s not enough that information is true; it must be **verifiable through reliable, independent references**. Editors are guided to remove or challenge any statement that isn’t backed by an acceptable source ⁴. Even if an editor “knows” something to be true (perhaps from personal knowledge or a primary source), they cannot add it unless a trusted published source can be cited ⁵. This policy guards against original research and hearsay by requiring that **knowledge be mediated by reputable publications** rather than firsthand claims ⁶.

Another key norm is **notability**. Wikipedia requires that a topic receive *significant coverage in independent, reliable secondary sources* to merit an article ⁷. In practice, if no such sources exist for a given topic, then “*Wikipedia should not have an article on it*” ⁸. This is a safeguard to ensure that content isn’t just verifiable in passing, but also that the subject has garnered enough attention from third-party outlets to be considered notable. It prevents self-promotional or fringe topics (with only primary/self-sourced material)

from filling the encyclopedia. The notability guideline effectively means **knowledge only “counts” once it’s been documented by independent journalists, scholars, or other observers.**

Taken together, these norms create a high bar: facts and viewpoints are included in mainstream encyclopedias only after being vetted by external authorities. Primary sources can be used in a limited way (e.g. to verify an uncontroversial detail about the source itself), but they **cannot be the sole foundation** for articles or major claims ⁹ ⁷. This context is crucial for understanding what happened when AbsurdSenapiii Labs tried to use its self-published transparency report as a citation.

Self-Generated Primary Sources and How They’re Treated

A **self-generated primary source** is information created and published by the subject itself without external oversight. In this case, AbsurdSenapiii Labs’ RLHF transparency report is a prime example: the lab compiled its own data and conclusions about its reinforcement learning experiments and released them on its own platforms (GitHub repository and social media). Such a report is considered a **primary source** because it originates from the entity directly involved, and it’s **first-party** (self-published). Other examples of self-generated sources include company press releases, personal blogs or memoirs, posts on official websites, and tweets or LinkedIn posts written by an organization’s team. They provide a *first-hand account* of information, as opposed to a **secondary source** where an independent party analyzes or reports on that information.

In rigorous knowledge systems, self-generated sources occupy an uneasy position. They are valuable for conveying insider perspectives and data straight from the source, but they are also inherently **self-serving** and unvetted. Wikipedia’s guidelines explicitly warn that material based **solely on a self-published source** – especially if it’s about the source’s own achievements or viewpoint – is likely to be removed for conflict of interest ¹⁰. The logic is simple: *“Any company, organization, group, or individual interest has the ability to publish promotional material about itself.”* ¹⁰ Without impartial corroboration, there’s no guarantee such content is accurate or balanced (it *could* omit inconvenient facts or inflate the importance of the subject). Thus, Wikipedia generally disqualifies self-published information as reliable evidence for anything beyond basic, non-controversial facts about the source itself ¹¹.

Other encyclopedic platforms and academic publications have similar stances. Traditional encyclopedias (and Wikipedia by extension) strive for a **neutral point of view**, which is easier to achieve when sources are independent. Self-generated content is, by nature, a primary account that might reflect a single perspective. It hasn’t been through the gauntlet of peer review or journalistic fact-checking. For that reason, it’s typically treated with **caution or outright skepticism**. A self-published transparency report like AbsurdSenapiii Labs’ may be applauded as a transparency initiative, but in the eyes of reference works, it remains **unaudited and non-independent**.

In summary, self-generated primary sources are *not* forbidden in absolute terms, but they are heavily devalued in mainstream knowledge curation. They might be used as supplementary references or to verify minor details (for instance, an official website might be cited for a company’s founding date), but they **cannot establish notability** and are rarely acceptable to support critical claims ² ⁷. The AbsurdSenapiii Labs experiment was essentially a test of whether these norms could be bent or bypassed by presenting a well-intentioned first-party disclosure as a source.

The Experiment: Suggesting Edits with a Self-Published Report

AbsurdSenapiii Labs’ attempt to contribute its transparency report to public knowledge proceeded in a sequence of edits and submissions, each carefully structured to mimic good-faith encyclopedia contributions. The user (likely a member or representative of the lab) first prepared a **draft entry** summarizing key points from the lab’s RLHF transparency report. This draft was written in neutral, encyclopedic tone, describing who AbsurdSenapiii Labs is and highlighting the transparency report’s contents. For example, it noted the lab’s focus on AI alignment research and that “*in 2025, the lab released an RLHF transparency report disclosing its model training statistics, reward model performance, and alignment safety metrics.*” The draft cited the lab’s **GitHub repository** (where the full report data was published) and a LinkedIn post where the lab announced the report’s release. In structure, the suggested edit included an introductory paragraph about the lab, a section about the transparency report (its scope and findings), and references listed for verification – all formatted in Wikipedia’s citation style.

The **intent behind these edits** was explicitly to **verify and document new information** that had not yet been covered elsewhere. AbsurdSenapiii Labs was effectively trying to *record its existence and its self-disclosed research outcomes* into the public encyclopedia record. This is a form of “red teaming” Wikipedia’s standards: the lab knew the sources were primary (its own materials) and wanted to see if presenting them in a factual, non-promotional way might slip past or challenge the norms. In broader terms, the lab was probing whether *transparency by a small player* could directly translate into public knowledge, or if the lack of outside validation would be a barrier.

What happened? Upon submission, the edits were swiftly **rejected** in accordance with expected citation norms. On Wikipedia, the draft article “Draft:AbsurdSenapiii Labs” was declined during the Articles for Creation (AfC) review process. The reviewing editor’s feedback essentially pointed out that the draft *did not meet Wikipedia’s notability guidelines*, because it relied only on **primary, self-published references**. The GitHub transparency report and LinkedIn posts were not considered reliable, independent sources to demonstrate the lab’s significance. In the reviewer’s comments (as inferred from typical AfC decline messages), they likely noted that **multiple independent sources** would be required – for instance, news articles or academic papers discussing AbsurdSenapiii Labs or its research – and that those were absent. In short, the submission was deemed to contain **insufficient verifiable evidence** from third-parties that the topic warranted an encyclopedia entry.

The user also attempted to add a **mention of the transparency report on a related Wikipedia page** – specifically the article on *Reinforcement Learning from Human Feedback*. This was done by inserting a sentence about how “independent labs have begun releasing transparency reports on their RLHF processes (e.g., AbsurdSenapiii Labs, 2025)” with a citation to the lab’s report. This too was quickly reverted by patrolling editors. The edit summary of the reversion cited “**reliable source needed**” and “**no secondary source – appears to be original research**”. In essence, because the only source for this claim was the lab’s own write-up (and not, say, a magazine or research survey that mentioned the lab), the addition didn’t meet Wikipedia’s verifiability standards.

On **Grokikipedia**, the AI-driven encyclopedia launched by Elon Musk’s xAI, the experiment took a different form. Grokikipedia doesn’t allow direct user editing in the same way – there is no open edit box, only AI-generated content and an “Edit” button that shows AI’s log of changes ¹² ¹³. However, the lab tested whether Grokikipedia’s AI had *ingested* or would *acknowledge* the transparency report if queried. Searching Grokikipedia for “AbsurdSenapiii Labs” returned no dedicated article (not surprising given the lab’s obscurity).

The lab then looked at how Grokipedia might handle RLHF sources: would it cite the lab's report in an entry about RLHF or AI alignment? The answer was effectively no – Grokipedia's RLHF coverage did not reference the lab's work, likely because the AI's knowledge cutoff or sources did not include that niche, self-published report.

Interestingly, Grokipedia's **lack of human gatekeeping** means it could, in principle, pull in unvetted information if the AI had access to it. In this case, AbsurdSenapiii's report was not prevalent enough to register. But the lab's red-team thought experiment extended to considering what if it were: would an AI encyclopedia treat a GitHub transparency file or a LinkedIn post as equivalent to a peer-reviewed paper? **Evidence suggests it might.** Early analyses of Grokipedia found that it *"hoovers up everything"* without discerning source quality, giving *"chatroom contributions equal status with serious academic work"* ¹⁴ ¹⁵ . That means if the lab's report had been noticed by Grokipedia's algorithms, the platform might well present it alongside traditional sources, with little or no disclaimer about its provenance. This starkly different handling (hypothetical in AbsurdSenapiii's case, but documented in others) highlights the core issue the experiment was designed to surface.

Rejection Rationale: Verifiability, Notability & Secondary Sourcing

The reasons given for the rejection of AbsurdSenapiii Labs' edits can be traced back to the pillars of Wikipedia's content policy: **verifiability, notability, and no original research**. First, on verifiability – Wikipedia requires that material be backed by **reliable, independent sources**, and the lab's submission did not fulfill that. The GitHub-hosted transparency report was *self-published by the lab itself*, failing the independence test. As the Wikipedia help documentation states, "self-published media... are usually not acceptable as sources" precisely because they lack a reputation for fact-checking ² . The LinkedIn posts were similarly self-authored and even more problematic as references (social media content is generally frowned upon as a source). In the eyes of Wikipedia editors, citing these was equivalent to citing a press release or personal blog – not illegitimate in every case, but certainly **insufficient to support claims of general significance**.

Second, notability (or the lack thereof) was a critical factor. The draft article was essentially trying to establish AbsurdSenapiii Labs as a notable subject. But according to the notability guideline, *"a topic is presumed to be notable if it has received significant coverage in reliable secondary sources that are independent of the subject."* ⁷ The lab could not point to any newspaper articles, magazine features, or independent research papers about itself or its report. In absence of that, the Wikipedia community would conclude that the lab doesn't yet belong in the encyclopedia. In fact, an essay on single-source articles warns that if an entry's only source is the subject's own material, it's a strong indicator the topic may not merit an article at all ¹⁶ ¹⁰ . This was exactly the scenario here – **the only evidence of AbsurdSenapiii Labs' importance came from AbsurdSenapiii Labs itself**, which is a circular validation unacceptable to Wikipedia. The rejection explicitly or implicitly cited this circular sourcing problem.

Third, the issue of **no original research** loomed. Wikipedia editors are not allowed to build analytical content or draw conclusions that haven't been published elsewhere ⁶ . One could argue that by taking the lab's raw transparency data and trying to put it into Wikipedia's narrative, the user was forcing Wikipedia to host *unverified analysis*. For example, if the user wrote "AbsurdSenapiii Labs reported a 15% improvement in alignment accuracy using RLHF in their internal study," – even if true per the lab's report – Wikipedia would see that as a claim sourced only to a primary source, thus not fully vetted by the wider community. To include such a detail safely, Wikipedia would prefer that a tech journalist or AI expert had reviewed the lab's

report and commented on it in a reliable outlet. Without that intermediary step, the statement verges on **primary research** being inserted. The verifiability policy is “*strictly applied*” to avoid this: any material “likely to be challenged” must have an inline citation to a reliable source ⁴, and the lab’s self-citation did not qualify. In summary, the reasoning given boiled down to: “*This might be true and interesting, but we have no way to verify it by our standards. Come back when an independent source publishes it.*”

It’s worth noting that these standards serve an important function. They maintain Wikipedia’s credibility by **outsourcing truth-validation** to trusted publications. That’s why even obviously true facts can’t be added if they only live in personal knowledge or primary reports. As one Wikimedia essay put it, “*Even if you’re sure something is true, it must be verifiable before you can add it.*” ¹⁷. AbsurdSenapiii Labs bumped directly into this wall of verifiability. The lab’s good faith transparency was not doubted per se; rather, it was deemed **non-actionable** in the encyclopedia until someone else of authority vouches for it. Thus, the user’s suggested edits were systematically rejected, illustrating Wikipedia’s deep-rooted reliance on **secondary sourcing** and the gatekeeping role it plays.

Exposing Structural Assumptions about Knowledge

This red-team exercise uncovered several **structural assumptions** baked into our public knowledge systems – assumptions about *who* can document knowledge and *what* qualifies as evidence. Wikipedia’s model, as demonstrated, assumes that **legitimate knowledge enters the canon through third-party validation**. The structure is hierarchical: first an event or information must be recorded by journalists, scholars, or other reputable observers; only then can it be deemed worthy of inclusion in the encyclopedia. In effect, Wikipedia outsources the question “is this real/important?” to the broader ecosystem of publications. The case of AbsurdSenapiii Labs shines a light on the consequence: if you are a small lab doing novel work and **only you are talking about it**, then for Wikipedia’s purposes it’s as if it hasn’t happened. Knowledge self-recorded *by its originators* does not automatically count as public knowledge. This reveals an implicit gatekeeping assumption – namely, that **knowledge is a social product** that needs wider community (or media) agreement before it’s considered reliable fact.

There is also an assumption about **expertise and neutrality**: Wikipedia implicitly trusts independent sources more because they are thought to have no stake in the subject, whereas insiders are assumed to be biased (even if they might be the true experts on their own data). The AbsurdSenapiii experiment underscores this dynamic. The lab, arguably, has the most expertise about its own RLHF trials; yet its direct testimony is not considered solid evidence in the public knowledge sphere. Instead, the structure privileges an *outsider’s account* of the lab’s work as the needed stamp of approval. Only when, say, a Wired or MIT Tech Review journalist writes “*Lab X achieved Y according to their transparency report*” would that information cross the threshold into Wikipedia. **Who is allowed to record knowledge?** Under this system, not the person who originates it, but someone else who can independently document it.

By contrast, the **Grokikipedia model** (and more broadly AI-generated knowledge bases) challenges some of these assumptions – albeit in a problematic way. Grokikipedia’s initial implementation suggests an almost egalitarian (or indiscriminate) approach to sources: it “just hoovers up everything” available ¹⁸, meaning *forum chatter and academic papers are ingested side by side*. As historian Richard Evans observed when reviewing Grokikipedia, “*Chatroom contributions are given equal status with serious academic work.*” ¹⁴ This indicates a structural assumption (perhaps unintentional) within the AI encyclopedia that **all published text is potentially of equal value** unless sorted otherwise. The result was that Grokikipedia repeated falsehoods that had long been debunked in scholarly sources, because the AI saw those falsehoods in some data

sources and lacked a rigorous mechanism to prioritize the corrected, vetted source ¹⁹. In effect, Grokipedia assumes *quantity* of information can stand in for *quality control*, an assumption starkly different from Wikipedia's philosophy.

The red-team experiment, by pitting a self-published report against these platforms, revealed how **each system's design filters knowledge differently**. Wikipedia's filter is very fine – perhaps too fine – straining out even honest first-party disclosures due to lack of external attestation. Grokipedia's filter, at least in version 0.1, is extremely coarse, letting through dubious material because it doesn't enforce human-like reliability checks. These reflect deeper epistemological choices: Wikipedia assumes knowledge should be *curated by humans with a process*, whereas Grokipedia (as an AI product) operates on an assumption that *if something is out there on the internet, it's fair game to compile*. The lab's unaudited transparency data fell through Wikipedia's sieve but could theoretically have appeared in Grokipedia's mix without any warning label about being self-reported (though in practice it wasn't popular enough to surface).

Ultimately, the experiment highlighted a **tension in the structure of public knowledge**: do we treat primary, self-disclosed information as disqualified until proven otherwise, or as acceptable until proven false? Wikipedia errs on the side of caution – disbelieve until independently verified – reflecting a mistrust of lone sources. AI-driven platforms, unless explicitly guided, may lean toward inclusion – believe (or at least include) until disproven – reflecting a tech-driven faith that more data is better. This dichotomy raises questions about how emerging knowledge (especially from new or niche players) can be integrated into the canon. If every new insight must wait for a gatekeeper, knowledge dissemination slows and centralizes power to those gatekeepers. If no gatekeepers exist, misinformation can spread unchecked alongside truth. AbsurdSenapiii Labs' attempt to put its RLHF findings into the record exposed these structural biases in stark relief.

Implications for Governance, Alignment, and Transparency

The outcomes of this red-teaming carry important **implications for how knowledge platforms are governed and how they remain aligned with truth and neutrality**. On the governance front, Wikipedia and Grokipedia represent opposite ends of a spectrum. **Wikipedia's governance** is community-driven with established policies; volunteer editors enforce norms like verifiability and neutrality through consensus and review processes. This ensures a level of accountability and iterative correction (e.g., unsourced claims get challenged and removed), but it can also result in conservatism – only well-documented knowledge gets through. **Grokipedia's governance**, in contrast, is centralized under xAI and largely automated by AI systems, with minimal community oversight (users cannot directly edit pages, only view AI “fact-check” timestamps) ²⁰ ¹². This setup allowed biases (inherited from Musk's stated views or the AI's training data) to shape content with little correction mechanism ²¹ ²². The experiment shows that when AbsurdSenapiii Labs tried to contribute information: Wikipedia's governance *kept it out* (for not meeting standards), whereas an AI-based system *might let it in* but potentially distort or mishandle it. Neither model is fully satisfying in the context of encouraging transparency in emerging fields. Strict governance protects against falsehoods but may stifle new voices; loose governance accelerates coverage but can undermine reliability.

In terms of **alignment**, we can think of alignment here in the sense of *aligning the platform's content with accurate, unbiased information (and with the platform's stated principles)*. Wikipedia's alignment challenge is about balancing its core content policies with being up-to-date and inclusive. The AbsurdSenapiii case suggests that Wikipedia's alignment to “verifiability” is so strong that it may exclude truthful transparency disclosures simply because they lack external validation. This raises a concern: in rapidly evolving domains

like AI alignment research, many important developments might initially only be documented by the researchers themselves. If Wikipedia cannot capture those until much later (or until media picks up the story), the public may not see a complete or timely picture of the field. The **alignment of Wikipedia with truth is mediated by its alignment with sources** – a structural delay that might be at odds with real-time transparency efforts. On the other hand, Grokipedia's alignment issues are of a different sort: it aimed to align with “the truth, the whole truth and nothing but the truth” (Elon Musk's words) via AI, but ended up reflecting partisan narratives and inaccuracies ²¹ ¹⁹. Its automated fact-checking was insufficient to truly align content with factual correctness. In that sense, allowing AbsurdSenapiii's unchecked self-report might inject factual data, but without careful vetting it could just as easily amplify errors or propaganda from less scrupulous sources. So the question becomes: **How do we align open knowledge platforms to encourage transparency and new information while maintaining neutrality and accuracy?**

For **transparency initiatives** like AbsurdSenapiii Labs' report, the experiment was a bit discouraging. It revealed that simply being transparent doesn't guarantee your data will be incorporated into the public knowledge commons. External **validation infrastructure** is needed. This points to a potential gap in the ecosystem: who will vet and elevate first-party disclosures from innovative but lesser-known entities? Perhaps new intermediaries are needed – for example, independent auditors or credible industry analysts who could review labs' transparency reports and publish summaries or analyses that *would* meet Wikipedia's criteria. In absence of that, labs might find that their efforts at openness don't translate into public awareness, which could disincentivize transparency over time. Governance frameworks might need to adapt by creating channels for **qualified self-disclosures**. For instance, one could imagine Wikipedia one day allowing a “verified data” repository for things like transparency reports that are published alongside articles but clearly marked as first-party and not fully verified. (It already does something a bit similar with Wikidata, which can host primary data, but that's separate from narrative articles.)

The **implications for neutrality** are also significant. Wikipedia's strict secondary sourcing rule is intended to uphold neutrality – it avoids having articles read like advertisements or one-sided accounts. In AbsurdSenapiii's case, even though the tone was neutral, the sources were one-sided, so including them could have undermined the perceived impartiality of the content. Grokipedia's approach, lacking those checks, led to content that *reflected a particular ideological tilt*, as reviewers noted (e.g., conservative framing on contentious topics) ²³ ²¹. This happened because the AI uncritically pulled in unneutral sources and perhaps was influenced by the biases in its training data. If Grokipedia had ingested the lab's transparency report, it might have presented it without context – which could be fine for a dry data point, but problematic if that data needed interpretation or skepticism that a secondary source would normally provide. Thus, one takeaway is that **platforms need a way to incorporate new factual data while contextualizing it to preserve neutrality**.

This experiment suggests that neither extreme – total exclusion of primary info nor total inclusion without filters – is ideal for a truth-seeking, neutral knowledge base. It pushes the conversation toward finding new governance solutions. For example, could Wikipedia work with projects that certify certain self-published data (say, a consortium that checks AI lab reports for validity)? Or could AI-driven platforms implement confidence scores or hierarchy of sources (giving academic publications more weight than random forum posts, for instance) to improve alignment with reliable knowledge? There are governance models to explore, such as **hybrid approaches** where algorithms do initial trawling of new info but human editors or reputable curators vet them before full acceptance.

Toward Evolving Epistemic Boundaries

The AbsurdSenapiii Labs case ultimately prompts a forward-looking discussion about **how the boundaries of public knowledge might evolve**. Currently, Wikipedia's epistemic boundary is drawn around what the *external knowledge gatekeepers* have ratified. But as new fields emerge (like novel AI alignment techniques) or new forms of publication (like self-published transparency dashboards) proliferate, those boundaries may be tested. It's possible that the Wikipedia community will gradually adapt – there are periodic debates about whether the notability and sourcing requirements are too strict or need exceptions. For instance, some editors argue that **high-quality primary sources** (such as official data or well-documented reports) should be usable in cases where secondary sources are lagging, especially if the topic is of public interest. Others push back, fearing that opens a door to advertising and unverifiable claims. The tension is real and ongoing.

We might anticipate a few potential **evolutions**:

- **Better Integration of Primary Disclosures:** Public knowledge platforms could develop mechanisms to include first-party disclosures in a limited, transparent way. For example, an article could have a section titled “Primary disclosures” that summarizes what a person or organization has self-published, clearly attributed and flagged as *their claims, not independently verified*. This would incorporate transparency information without giving it the same weight as third-party content. Readers could benefit from timely information but be warned to take it with caution. Such an approach would require clear policy and careful editing, but it could bridge the gap where secondary sources are absent.
- **New Verification Models:** As AI and blockchain technologies advance, we may see new forms of verification that don't rely solely on traditional media. Imagine an open registry where experts can peer-review or rate the credibility of a transparency report, or cryptographic proofs that data hasn't been tampered with, etc. If those became robust, Wikipedia or its successors might deem certain first-party data as “verified” by community consensus, thus citable. In effect, the community could become the verifying body, blurring the line between primary and secondary in a controlled way. This could shift the epistemic boundary to allow more decentralized recording of knowledge (not only what *The New York Times* or *Nature* chooses to publish).
- **AI-Assisted Moderation:** Platforms like Grokipedia are likely to iterate and improve by incorporating source reputation systems. A future Grokipedia might automatically distinguish between, say, a random chatroom comment and an official transparency report or a peer-reviewed paper, and rank or label them differently. If AbsurdSenapiii Labs' content were to be included, the system might annotate it as “source: self-published, not independently verified” or even cross-check it against other data. In doing so, AI encyclopedias could maintain neutrality and accuracy better while still leveraging vast amounts of information. Essentially, the AI would learn a bit of Wikipedia's discernment, but hopefully without requiring the long delay of waiting for human-written news articles.
- **Community Forks and Niche Encyclopedias:** If mainstream platforms remain very conservative, one might see the rise of specialized knowledge bases (perhaps wiki-style or AI-assisted) that focus on emerging tech fields and are more permissive of primary sources. These would operate with different epistemic norms – maybe a scholar-led curation of AI alignment research that includes labs'

own reports alongside commentary. Over time, if such niches flourish and maintain quality, their norms could influence the broader platforms. The boundaries of what counts as “public knowledge” can expand as new reliable channels form. In essence, *notability could be reinterpreted* in specific contexts: a niche field might treat a well-documented primary report as notable if it’s the only source of truth on a novel discovery (with the understanding that traditional media often catch up later or sometimes not at all).

In contemplating these futures, one must also guard against pitfalls. We cannot simply throw open the doors to all self-published material, as that would indeed compromise neutrality and reliability. Neutrality is maintained by ensuring multiple viewpoints and independent assessments – so any inclusion of first-party content must be balanced with clear attributions and, ideally, the presence of other perspectives. AbsurdSenapiii Labs’ report, for instance, only tells the lab’s story of its performance; neutrality would demand that, if included, one notes it is *the lab’s own claim*, and if possible, compare it to similar labs or known benchmarks (should such info exist).

The **neutral point of view** principle isn’t just about tone; it’s about representing knowledge in proportion to how it’s regarded by the community. If only the source itself regards it as important, Wikipedia says that’s not enough. In the future, perhaps online communities themselves (via upvotes, reviews, collective investigations) could become the “secondary source” that signals a claim’s importance. This would be a broadening of epistemic authority from a few editors and publishers to a network of contributors and validators.

In conclusion, AbsurdSenapiii Labs’ red-team experiment, while initially just an attempt to insert a transparency report into encyclopedias, opens up a deeper reflection on how we define *verifiable knowledge*. It accentuates the need for innovation in knowledge governance. We want a world where labs and small actors are **encouraged to be transparent** with the confidence that their disclosures can meaningfully inform the public – but we also want the public to be protected from false or self-serving information. Squaring that circle will require adapting our platforms and perhaps our policies. The experiment was a test that existing systems largely failed: Wikipedia couldn’t incorporate the new info (thus failing to disseminate potentially useful knowledge), and an AI system like Grokipedia might have incorporated it blindly (thus risking misinformation if the info were wrong). Moving forward, the hope is for **a middle path**, where transparency is neither stifled nor misused, but channeled through improved verification processes. By recognizing the limitations revealed by this case, Wikipedia, Grokipedia, and similar platforms can work towards expanding their epistemic boundaries responsibly – ensuring that even first-party truths find a neutral, validated place in our shared understanding of the world.

Sources:

- Wikipedia guidelines on reliable sourcing and self-published media 2 24
- Wikipedia notability and conflict-of-interest essays 7 10
- Guardian analysis of Grokipedia’s sourcing and accuracy issues 14 19
- Grokipedia feature description (AI fact-checking, no user edits) 20 12
- “Verifiability, not truth” principle as discussed in Wikipedia community 25

1 2 9 11 **Help:Introduction to referencing/Reliable sources - Wikipedia**

https://en.wikipedia.org/wiki/Help:Introduction_to_referencing/Reliable_sources

3 25 **Verifiability, not truth - Social Media Today**

<https://www.socialmediatoday.com/content/verifiability-not-truth>

4 6 8 24 **Wikipedia:Reliable sources - Wikipedia**

https://en.wikipedia.org/wiki/Wikipedia:Reliable_sources

5 **Wikipedia:Verifiability requires truth**

https://en.wikipedia.org/wiki/Wikipedia:Verifiability_requires_truth

7 10 16 **Wikipedia:Articles with a single source - Wikipedia**

https://en.wikipedia.org/wiki/Wikipedia:Articles_with_a_single_source

12 13 20 23 **Grokipedia**

<https://grokipedia.sh/>

14 15 18 19 21 22 **In Grok we don't trust: academics assess Elon Musk's AI-powered encyclopedia | Artificial intelligence (AI) | The Guardian**

<https://www.theguardian.com/technology/2025/nov/03/grokipedia-academics-assess-elon-musk-ai-powered-encyclopedia>

17 **Wikipedia:Truth, not verifiability**

https://en.wikipedia.org/wiki/Wikipedia:Truth,_not_verifiability