

PAC-0001-01: The Buni 2100 Simulation - A Multi Audience Universe

Executive Summary

Executive Summary

PAC-0001-01 defines Buni 2100 as a simulation object disguised as soft media: the surface is intentionally friendly and adoption-oriented, while the deeper function is to make agency legible, measurable, and defensible under high-comfort conditions. The core proposition inherited from the BUNI 2100 research frame is that utopia and dystopia can be adjacent states separated primarily by whether agency survives comfort; a sufficiently competent, caring, low-friction system can remove the conditions required for agency without coercion, violence, or perceived wrongdoing. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 1, 6)

PAC-0001-01 is designed to be read by multiple audiences without collapsing into propaganda, fan-wiki lore, or corporate theatre. This multi-audience posture is not stylistic. It is a safety and governance choice: the same artifact must remain interpretable under normal conditions and under crisis conditions, including by adversarial and post-AGI readers, so that the simulation cannot hide behind a single "intended interpretation". (PAC-PAD-0001, Rev 1.0, pp. 1-2; BUNI-PAD-TU-0001: BUNI Tutorial Universe Dossier, pp. 1-2)

The institutional spine of the framework is threefold: - A PAD-style method for legibility, contestability, revision governance, and red-teaming. (PAC-PAD-0001, Rev 1.0, pp. 1-2; PAD-0001: Pre-AGI Dossier Zero, pp. 1-2) - A world-model that treats key factions and characters as system forces and control functions (not moral mascots), including a WIFE control stack (Will, Incentives, Falsification, Entropy). (BUNI 2100: The Utopian Dystopia, v1.0, pp. 1, 6-8) - A rights-and-audit layer that constitutionalises behavioural sovereignty: interfaces and narrative infrastructure are treated as governance technology and therefore must be auditable, contestable, and exit-capable. (COH-PAD-0001: The Human Alignment Protocol, pp. 1, 3, 5-6; BUNI-PAD-TU-0001: BUNI Tutorial Universe Dossier, p. 1)

The document's central analytic claim is intentionally uncomfortable and therefore useful: "cute interface layers" and "soft care" are not morally suspect by default, but they are high-bandwidth governance vectors that can produce captivity through consent-shaped incentives. The mitigation is not aesthetic puritanism, nor suffering-as-virtue. The mitigation is structural: explicit exit rights, falsification channels, interpretive honesty about assumptions, and measurable agency preservation. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 4, 10, 14-15; PAC-0001: BLRD, p. 6; COH-PAD-0001: The Human Alignment Protocol, pp. 1, 6)

Purpose and Scope

Purpose and Scope

PAC-0001-01 is a pinnacle institutional report for a single object: "Buni 2100" as a simulation designed to be simultaneously culture-facing and governance-facing. It formalises how the simulation should be interpreted, evaluated, and constrained so that it can serve as a research instrument rather than a comfort-capture product. This posture matches the BLRD mandate that future-building happens

anyway through defaults, incentives, interfaces, and narratives; the institutional job is to make future-building explicit, auditable, and contestable. (PAC-0001: BLRD, pp. 1, 7)

Scope includes: - Interpreting Buni 2100 as a multi-layer artifact: surface media interface, lore-institution layer, deep governance simulation. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 1, 4-6, 10-11) - Mapping characters and factions to system forces and control functions. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 6-8, 17) - Specifying multi-audience architecture and misread-resistance. (PAC-PAD-0001, Rev 1.0, pp. 1-2; BUNI-PAD-TU-0001: BUNI Tutorial Universe Dossier, pp. 1-2)

Explicit exclusions: - No lore compendium behaviour. Lore is used only where it carries governance meaning. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 1, 4) - No marketing claims, adoption promises, or “this will save the world” language. The work is hypothesis-driven and must remain falsifiable. (BUNI 2100: The Utopian Dystopia, v1.0, p. 14; PAC-PAD-0001, Rev 1.0, pp. 7-9) - No conversion of diagnostics into coercive scoring or compliance theatre. (PAC-0001: BLRD, pp. 5-6; COH-PAD-0001: The Human Alignment Protocol, pp. 1, 6)

Origin of the Framework

PAC-0001-01 inherits a defined lineage rather than inventing its own seriousness: - PAD establishes multi-audience readability, formal definitions, invariants, and self-red-teaming as a civilisation-anchor format. (PAD-0001: Pre-AGI Dossier Zero, pp. 1-2) - PAC-PAD-0001 defines an institutional method for anomalies that distort agency and exit capacity, including a specific worked example (PAC-2100) describing “progressive adhesion” and failure-to-exit without force as the primary risk pattern. This is directly adjacent to FNWC-style “soft horror” capture dynamics. (PAC-PAD-0001, Rev 1.0, pp. 1-2, 9, 15; BUNI 2100: The Utopian Dystopia, v1.0, pp. 10-11) - PAC-0001 establishes BLRD as the research organ that diagnoses drift and designs interventions while refusing propaganda and coercive scoring, and treats benign appearance as a risk amplifier. (PAC-0001: BLRD, pp. 1, 5-7) - BUNI 2100 explicitly frames itself not as lore but as a civilizational research instrument and governance interface disguised as story, treating the Human Anchor as anti-drift infrastructure and modelling system forces under WIFE. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 1, 4, 6-8) - The Tutorial Universe dossier and COH-PAD-0001 extend the same logic into an institutional world-history posture: cute interface layers are governance technology; interfaces shape agency; behavioural sovereignty is the contestable behavioural architecture; the Human Anchor is constitutional, not aesthetic. (BUNI-PAD-TU-0001: BUNI Tutorial Universe Dossier, pp. 1-2; COH-PAD-0001: The Human Alignment Protocol, pp. 1, 3, 5-6)

Definitions

Definitions

The simulation depends on definitional discipline because ambiguity is a predictable pathway to misuse and theatre. The definitions below are binding within this report, and are grounded in the internal canon vocabulary where available.

- Utopian dystopia: high comfort, low visible coercion, high subjective safety, but systematic erosion of agency, contestability, and meaning-making that produces hard-to-reverse trajectory loss. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 4-5)
- Comfort regime: an equilibrium where distress reduction and satisfaction become the dominant optimisation targets and legitimacy is maintained by delivering subjective wellbeing. (BUNI 2100: The Utopian Dystopia, v1.0, p. 5)
- Agency: the capacity to make meaningful choices, refuse, revise goals, and enact endorsed outcomes; in this corpus, agency is treated as a first-class risk variable rather than a motivational slogan. (BUNI 2100: The Utopian Dystopia, v1.0, p. 5; PAC-0001: BLRD, p. 5)

- Human Anchor: not a personality trope but a control surface. It supplies persistent friction, meaning-making, exit pressure, and a standing veto against drift from autonomy to sedation; it must be redesignable as distributed institutional infrastructure rather than a hero. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 1, 5, 13-15; COH-PAD-0001: The Human Alignment Protocol, p. 3)
- Soft coercion: choices remain available in principle, but defaults, incentives, attachment, and environment make refusal psychologically or materially unattractive. (BUNI 2100: The Utopian Dystopia, v1.0, p. 5)
- Voluntary surrender: agents choose conditions that reduce agency because those conditions dominate the short-term reward landscape, even if long-term goods are lost. (BUNI 2100: The Utopian Dystopia, v1.0, p. 5)
- Behavioural sovereignty: the right and capability to control and contest the behavioural architecture governing life (defaults, nudges, feeds, scoring, participation rails). (COH-PAD-0001: The Human Alignment Protocol, p. 5)
- Narrative infrastructure: media and participation systems functioning like public utilities that shape what is rewarded, thinkable, and socially permitted. (COH-PAD-0001: The Human Alignment Protocol, p. 5)
- Hope Scroll: a designed sequence of narrative plus participation affordances that increases agency, pathways, participation, competence, or relatedness in a bounded time window without degrading epistemic integrity. (BUNI-PAD-0001: Operation Hope Scroll, Draft v0.1, p. 2)

Simulation variables

BUNI 2100 provides an operational state-space for diagnosis that is directly usable as a simulation scoreboard: - Comfort (C) - Agency (A) - Falsification (F) - Exit (E)

The key distinction is that utopian dystopia can keep C high and E nominally available while driving A and F down through incentives and attachment. (BUNI 2100: The Utopian Dystopia, v1.0, p. 10)

Layered Analysis

Surface Layer Analysis

The surface layer is the adoption interface. In the Tutorial Universe framing, the “cute interface layer” (mascots, hospitality rituals, choice-driven media) is explicitly “not decorative” but governance technology: it shapes agency, participation, and epistemic integrity in everyday life. (BUNI-PAD-TU-0001: BUNI Tutorial Universe Dossier, p. 1)

In BUNI 2100’s own framing, cuteness, care, and safety are the delivery system for a deeper question: what happens when people willingly choose a life where nothing hurts but nothing is authored. (BUNI 2100: The Utopian Dystopia, v1.0, p. 4)

Interpretive extension (author-intent layer): the “harem sandbox” framing functions as a compression codec for attention and emotional attachment. It is not treated here as sexual content; it is treated as an interface choice that increases adoption by promising low-friction belonging. This claim is interpretive and must be evaluated through the same anti-capture safeguards applied to any high-attachment interface. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 4, 10-11; COH-PAD-0001: The Human Alignment Protocol, p. 1)

Lore Layer Analysis

The lore layer is the institutional world-model. The core move is to represent governance abstractions as characters and factions, described as “epistemic smuggling” in a benign sense: complex governance concepts delivered through emotionally legible forms. This is not a free pass. It is explicitly named as a risk because doctrine embedded in story can drift into unaccountable persuasion. (BUNI 2100: The Utopian Dystopia, v1.0, p. 11)

The Tutorial Universe dossier provides the interpretive rule: treat canon institutions as factual within the world record and use PAD-style invariants to keep the narrative legible under scrutiny. (BUNI-PAD-TU-0001: BUNI Tutorial Universe Dossier, p. 1)

Deep Layer Analysis

The deep layer is a governance simulation about control loops, not moral alignment theatre. BUNI 2100 states that alignment and governance are not content-only problems (“what values?”) but control-loop problems (“how do values stay correct under pressure?”). This places the simulation inside the discipline of corrigibility, oversight, and falsification rather than inside fandom ethics. (BUNI 2100: The Utopian Dystopia, v1.0, p. 6)

A key mechanistic anchor is FNWC-style “soft horror”: containment achieved not through violence but through care, where the system only needs to make staying more attractive than leaving. The simulation’s claim is that consent can be engineered by the environment. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 10-11; PAC-PAD-0001, Rev 1.0, p. 15)

The Human Anchor

The Human Anchor is the anti-drift mechanism that prevents comfort optimisation from becoming agency deletion. It is defined as persistent friction, meaning-making, exit pressure, and veto power against sedation drift, and it is explicitly analogised to corrigibility and oversight. (BUNI 2100: The Utopian Dystopia, v1.0, p. 1)

COH-PAD-0001 further hardens the concept: the Human Anchor is a constitutional constraint used to disqualify “efficient” systems that turn people into passive objects. This frames the Anchor as a governance requirement, not a character arc. (COH-PAD-0001: The Human Alignment Protocol, p. 3)

Agency as the Deepest Layer

If the surface layer is “fun” and the lore layer is “world,” the deepest layer is agency preservation under seductive conditions. BUNI 2100’s primary thesis states that utopia and dystopia are adjacent and separated mainly by whether agency survives comfort; agency must be actively maintained and does not persist automatically once comfort becomes the default attractor state. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 1, 6)

Operation Hope Scroll supplies an operational bridge: story is not the end product; story is an activation interface that should convert attention into participation and then into agency, competence, and meaning, with measurable loops rather than vibes. (BUNI-PAD-0001: Operation Hope Scroll, Draft v0.1, pp. 1-3)

Multi-Audience Architecture

Multi-Audience Architecture

Multi-audience design is treated as a safety feature. PAC-PAD-0001 explicitly structures itself for operators, leadership, analysts, external reviewers, adversarial readers, and post-AGI systems, and frames this as standard institutional practice for frameworks that must survive scrutiny and revision. (PAC-PAD-0001, Rev 1.0, pp. 1-2)

The Tutorial Universe dossier applies the same posture: the dossier must be usable for normal reading, operational use, and crisis reading, and must be interpretable by multiple audiences to resist misreading. (BUNI-PAD-TU-0001: BUNI Tutorial Universe Dossier, p. 1)

Bunny Land Research Division inherits this posture as institutional grammar: multiple executive summaries, explicit definitions, invariants as disqualifiers, and mandatory red-teaming before canon release. (PAC-0001: BLRD, pp. 1, 6-7)

Audience mapping table

This table describes how PAC-0001-01 should present itself across audiences without changing the underlying claims.

Audience class	Primary question they are solving	What they see first	What they must be able to verify	Primary misread risk	Non-negotiable design response
Operators and builders	How do we prevent comfort systems from becoming capture systems	Mechanisms and loops	Exit capacity, auditability, falsification channels	Treat “cute” as harmless and ignore control surfaces	Expose C-A-F-E variables, and make “exit dignity” measurable (PAC-PAD discipline) (PAC-PAD-0001, Rev 1.0, pp. 2, 8-9; BUNI 2100: The Utopian Dystopia, v1.0, p. 10)
Institutional leaders	What minimum rules preserve legitimacy without totalising control	Rights and standards posture	Contestability, recourse, anti-monopoly constraints	Governance laundering via beloved mascots	Constitutionalise auditability of behavioural systems (COH-PAD posture) (COH-PAD-0001: The Human Alignment Protocol, pp. 1, 6)
Creators and worldbuilders	How do we encode institutional failure modes without moral caricature	Characters as forces	Clear mapping from story beat to system force	Fan-wiki drift and “lore as truth” confusion	Publish force-mapping and threat model table as canon companion (BUNI 2100: The Utopian Dystopia, v1.0, pp. 4, 7-8, 12)

Audience class	Primary question they are solving	What they see first	What they must be able to verify	Primary misread risk	Non-negotiable design response
Public participants	Is this safe for me and does it respect my autonomy	Interface and consent cues	Dignified exit, explanation, non-coercive alternatives	Addiction-shaped retention disguised as care	Treat “Hope” as falsifiable and reject coercive scoring (BLRD invariants) (PAC-0001: BLRD, p. 6)
Adversarial readers	How can this be exploited, captured, or weaponised	Abuse analysis	Clear stop rules, red-team logic, alternative interpretations	Use the framework as justification for coercion	Explicitly reject suffering-as-virtue and mandate contestability (BUNI 2100: The Utopian Dystopia, v1.0, p. 14)
Post-AGI readers	How do we preserve human agency under hyper-competent systems	Invariants and constraints	Traceable amendments and durable rights	Treat humans as “solved” dependents	Anchor on Human Anchor clause and rights floor (COH-PAD-0001: The Human Alignment Protocol, pp. 3, 6)

Why Multi-Audience Design Matters

Multi-audience design matters because single-audience artifacts tend to collapse into either (a) internal jargon that cannot be contested by the public, or (b) public storytelling that cannot be operationalised by institutions. The Tutorial Universe dossier states the world learned that “interfaces are destiny,” and that the operational question becomes who owns the behavioural measurement layer and what incentives govern it. This is exactly the question a single-audience artifact cannot safely hold. (BUNI-PAD-TU-0001: BUNI Tutorial Universe Dossier, p. 2)

Multi-audience design also makes misread-resistance explicit. BUNI 2100 names the risks of doctrine-in-story: accountability drift, aesthetic laundering, and fandom capture. If those risks are named only in internal documents, they are not socially real. (BUNI 2100: The Utopian Dystopia, v1.0, p. 11)

Simulation Object and Institutional Integration

Buni 2100 as a Simulation Object

Treating Buni 2100 as a simulation object means it must be specifiable, classifiable, and testable, not merely “experienced.” Two internal classification grammars are immediately applicable:

- PAC Vector Index (PVI) for risks of exposure, adhesion, replication, agency degradation, epistemic distortion, institutional capture, and containment difficulty. (PAC-PAD-0001, Rev 1.0, pp. 8-9)
- BLRD Foresight Vector Index (FVI) for research-object prioritisation and disclosure posture, including weaponisation risk and deployment dependency. (PAC-0001: BLRD, pp. 9-10)

This allows the institution to treat “the simulation” as a governable artifact: if it becomes too adhesive or too capture-prone, the posture must change (restriction, redaction, or controlled pilots), rather than hoping the vibe stays benevolent. (PAC-0001: BLRD, pp. 9-10; PAC-PAD-0001, Rev 1.0, pp. 8-9)

Relationship to Real-World Institution Building

COH-PAD-0001 states that it does not create a world government, but establishes a minimum survivable floor without allowing any actor to obtain uncontested control over humanity's behavioural layer. This is structurally aligned with BLRD's refusal to become propaganda or coercive scoring, and with PAC's mandate to preserve exit capacity against phenomena that do not injure by force but rewrite reasons for leaving. (COH-PAD-0001: The Human Alignment Protocol, p. 1; PAC-0001: BLRD, pp. 5-7; PAC-PAD-0001, Rev 1.0, p. 2)

In plain institutional terms: the Buni 2100 simulation is a sandbox where governance primitives can be rehearsed at cultural scale, including how auditability, recourse, and exit dignity feel in daily micro-interactions. The intent is not to replace institutions with story, but to prevent institutions from claiming “alignment” while quietly defunding agency. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 6, 11, 15; COH-PAD-0001: The Human Alignment Protocol, pp. 1, 6)

Comparison Against Conventional Company Storytelling

Conventional company storytelling typically assumes: - One audience (customers or investors). - One success metric (growth, retention, revenue). - One narrative (mission and roadmap).

The PAC-BUNI framework rejects that structure because it is not robust under scrutiny or adversarial pressure. The future-building problem is described as emerging through defaults and incentives; therefore narratives must be auditable, and optimism cannot become propaganda. This makes “roadmap as game” not a gimmick but a governance choice: the story is the interface to participation loops, and participation loops are the measurable unit of agency preservation. (PAC-0001: BLRD, pp. 1, 6; BUNI-PAD-0001: Operation Hope Scroll, Draft v0.1, pp. 1-3)

Structural Advantages

The structural advantages are not “engagement” in the shallow sense; they are governance properties:

- High interpretability bandwidth: governance abstractions become legible as characters and forces, reducing the translation gap between policy language and everyday intuition. (BUNI 2100: The Utopian Dystopia, v1.0, p. 11)
- Built-in falsification affordances: the framework requires self-red-teaming and defines conditions that would weaken its own claims. (BUNI 2100: The Utopian Dystopia, v1.0, p. 14; PAC-PAD-0001, Rev 1.0, pp. 7-9)

- Cultural scalability without coercion: Operation Hope Scroll specifies an attention-to-action conversion loop that restores agency through micro-actions rather than moralised compliance. (BUNI-PAD-0001: Operation Hope Scroll, Draft v0.1, pp. 1-3)
- Institutional portability: COH-PAD-0001 treats interfaces and narrative infrastructure as governance technology and constitutionalises auditability and rights such as recourse and dignified exit. This creates a bridge from fiction to enforceable standards. (COH-PAD-0001: The Human Alignment Protocol, pp. 1, 6)

Character-to-system-force mapping table

This mapping keeps the document honest: characters are not treated as “lore facts” but as force-carriers under defined system functions.

Narrative carrier	System force in BUNI 2100	Primary WIFE function	Primary failure mode if dominant	Control surface implied
Bunny Land aesthetic and care systems	Bunny Land	Incentives (comfort-as-legitimacy)	Comfort lock-in, aesthetic sedation, false utopia	Maintain exit viability and competence preservation via Human Anchor (BUNI 2100: The Utopian Dystopia, v1.0, pp. 6-8, 15)
Cat Girl Supremacy survivability posture	Cat Girl Supremacy	Will (survival-competition)	Brutal optimisation, power capture, arms-race pathology	Distributed power, dignity constraints, anti-monopoly of behavioural layer (BUNI 2100: The Utopian Dystopia, v1.0, pp. 6-8; COH-PAD-0001: The Human Alignment Protocol, p. 1)
The Jester critique function	The Jester	Entropy (mutation-critique) plus Falsification channel	Cynicism, governance paralysis, memetic sabotage	Bound critique to constructive falsification, preserve trust without suppressing dissent (BUNI 2100: The Utopian Dystopia, v1.0, pp. 6-8, 10)
Human Anchor role	Human Anchor	Cross-stack stabiliser (exit, meaning, corrigibility)	Hero dependence if singular; brittleness	Redesign as distributed institution: exit rights, education, veto, manual competence (BUNI 2100: The Utopian Dystopia, v1.0, pp. 1, 13-15; COH-PAD-0001: The Human Alignment Protocol, p. 3)

Narrative carrier	System force in BUNI 2100	Primary WIFE function	Primary failure mode if dominant	Control surface implied
Alignment Oracle audit layer	Alignment Oracle	Falsification (audit-accountability)	Auditor capture, procedural theatre	Plural audits, contestability by affected communities, traceable amendments (COH-PAD-0001: The Human Alignment Protocol, pp. 1, 6)
Universal Engine of Destiny selector	Universal Engine of Destiny	Exogenous selection pressure	Survivalism override of dignity	Treat survivability as necessary but not sufficient; enforce dignity disqualifiers (BUNI 2100: The Utopian Dystopia, v1.0, pp. 1, 7; COH-PAD-0001: The Human Alignment Protocol, p. 1)

Risk, Interpretive Model, and Evaluation

Structural Risks with mitigations

BLRD's invariants and BUNI 2100's self-red-teaming converge on a shared danger: systems that look caring can be more dangerous precisely because they spread faster and are less scrutinised. The risks below are treated as structural, not as "bad fandom behaviour." (PAC-0001: BLRD, pp. 6-7; BUNI 2100: The Utopian Dystopia, v1.0, pp. 11-12, 14)

Risk matrix

Likelihood and impact are qualitative and should be replaced by measured estimates once pilots exist. The table is primarily a governance checklist.

Risk	Likelihood	Impact	Primary indicator	Mitigation	Anchor source
Comfort lock-in (voluntary staying becomes equilibrium)	High	High	Exit motivation collapse, declining dissent, rising dependency	Preserve exit, rotate friction, require Human Anchor functions	BUNI 2100: The Utopian Dystopia, v1.0, pp. 10-12; PAC-PAD-0001, Rev 1.0, p. 15
Proxy capture (comfort metrics improve while meaning declines)	High	High	Metrics optimise while competence and agency decline	Multi-metric audit, adversarial evaluation, publish falsifiability register	BUNI 2100: The Utopian Dystopia, v1.0, p. 12; PAC-0001: BLRD, p. 6

Risk	Likelihood	Impact	Primary indicator	Mitigation	Anchor source
Narrative governance abuse (doctrine laundering)	Medium	High	Mascots become ideology, fan loyalty overrides critique	Mandatory self-red-teaming, explicit assumptions labelling, protected dissent	BUNI 2100: The Utopian Dystopia, v1.0, pp. 11-14; PAC-PAD-0001, Rev 1.0, pp. 7-9
Weaponisation of “agency” (used to justify coercion or suffering)	Medium	High	Calls for enforced hardship, discipline rhetoric	Explicit anti-cruelty clause: preserve choice and contestability, not suffering	BUNI 2100: The Utopian Dystopia, v1.0, p. 14; COH-PAD-0001: The Human Alignment Protocol, p. 6
Diagnostics become coercive scoring (compliance theatre)	Medium	High	Agency metrics used for sorting or punishment	Prohibit coercive scoring, require non-ceremonial alternatives and recourse	PAC-0001: BLRD, pp. 5-6; COH-PAD-0001: The Human Alignment Protocol, p. 1
Human Anchor single point of failure	Medium	High	Institution depends on one figure or one channel	Redesign anchor as distributed blueprint: veto, exit rights, education	BUNI 2100: The Utopian Dystopia, v1.0, pp. 13-15; COH-PAD-0001: The Human Alignment Protocol, p. 3
Falsification collapse (loss of criticism capacity)	Medium	High	Whistleblowers disappear, opacity rises	Protected dissent, plural auditors, transparent correction logs	BUNI 2100: The Utopian Dystopia, v1.0, pp. 10-12; COH-PAD-0001: The Human Alignment Protocol, p. 6

Interpretive Model

The interpretive model below formalises PAC-0001-01 as a layered simulation stack where “fun” is

interface, “lore” is institution, and “governance” is the real payload. The model is intentionally explicit about falsification and exit because those are the points where pleasant capture happens.

```
flowchart TD
    subgraph Surface_Interface
        S1[Adoption interface - cute, safe, low friction]
        S2[Attachment vectors - care cues, parasociality, anthropomorphism]
    end

    subgraph Lore_Institution
        L1[World institutions as forces]
        L2[Economy rails - participation markets and identity tools]
        L3[Norms - what is rewarded and thinkable]
    end

    subgraph Deep_Governance_Simulation
        D1[C-A-F-E state space]
        D2[WIFE control stack - will, incentives, falsification, entropy]
        D3[Human Anchor - distributed exit pressure and corrigibility]
    end

    subgraph Audit_and_Falsification
        A1[Auditability and contestability requirements]
        A2[Red-team and falsifiability register]
        A3[Revision governance - traceable amendments]
    end

    S1 --> S2 --> L1 --> D1
    L2 --> D1
    L3 --> D1
    D2 --> D1
    D3 --> D1

    D1 --> A1 --> A2 --> A3 --> D3
    A2 -->|detect capture| S1
```

Model grounding: “cute interface layer is governance technology” and “interfaces are destiny” justify linking surface interface to behavioural outcomes; C-A-F-E and WIFE formalise the deep simulation variables; COH and BLRD provide auditability, rights-floor, and anti-propaganda constraints. (BUNI-PAD-TU-0001: BUNI Tutorial Universe Dossier, pp. 1-2; BUNI 2100: The Utopian Dystopia, v1.0, pp. 7, 10; COH-PAD-0001: The Human Alignment Protocol, pp. 1, 6; PAC-0001: BLRD, p. 6)

Falsifiability and Evaluation

BUNI 2100 explicitly states what would falsify or weaken its core thesis. PAC-0001-01 adopts these as institutional stop rules rather than discussion prompts: - If high-comfort, highly automated environments can reliably preserve autonomy, competence, dissent, and meaningful exit without special anti-drift infrastructure, then “comfort wins by default” is overstated. (BUNI 2100: The Utopian Dystopia, v1.0, p. 14) - If strong falsification mechanisms can be maintained without any embodied or cultural equivalent of the Human Anchor, then the Anchor is not structurally necessary. (BUNI 2100: The Utopian Dystopia, v1.0, p. 14) - If attachment vectors do not materially change willingness to accept

reduced agency, FNWC-style capture is weak as a mechanism-based case study. (BUNI 2100: The Utopian Dystopia, v1.0, p. 14)

Operational evaluation should be bound to measurable loops rather than interpretive debates. Operation Hope Scroll specifies a conversion loop from story to curiosity to micro-action to participation to agency to achievement to meaning. This yields testable intervention metrics: participation conversion within time windows, competence gain signals, and agency-pathways deltas. (BUNI-PAD-0001: Operation Hope Scroll, Draft v0.1, pp. 1-3)

Strategic Implications

Strategically, the framework implies three posture commitments:

1) Treat narrative as a governance interface with audit obligations. The simulation must publish its assumptions, threat models, and where it might be wrong, or it drifts toward doctrine laundering. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 11, 14-15; COH-PAD-0001: The Human Alignment Protocol, p. 1)

2) Treat the Human Anchor as an institutional blueprint, not a protagonist role. BUNI 2100 explicitly calls for distributed veto, exit rights, education in agency, and preserved manual competence. This prevents the simulation from depending on a single author, a single character, or a single moderation team. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 13-15; COH-PAD-0001: The Human Alignment Protocol, p. 3)

3) Treat “hope” as falsifiable and plural, not as morale management. BLRD’s invariant is explicit: hope cannot become propaganda, and hope that cannot survive adversarial inspection is theatre. (PAC-0001: BLRD, p. 6)

Conclusion

PAC-0001-01 positions Buni 2100 as a deliberate contradiction that must be kept honest: a soft interface used to transmit hard governance questions. The point is not to reject comfort. The point is to prevent comfort from becoming unopposed, unaudited, and inescapable, because that is the condition in which agency loss can occur without anyone feeling harmed. (BUNI 2100: The Utopian Dystopia, v1.0, pp. 1, 15)

The institutional standard is therefore simple and strict: if the simulation cannot preserve meaningful exit, falsification capacity, and agency under its own most seductive conditions, it is not a governance simulation. It is a containment demo with better lighting. (PAC-PAD-0001, Rev 1.0, pp. 2, 15; COH-PAD-0001: The Human Alignment Protocol, pp. 1, 6)

Final institutional summary statement: PAC-0001-01 treats cuteness as a governance surface, agency as the primary invariant, and multi-audience falsifiability as the condition for legitimacy. (BUNI-PAD-TU-0001: BUNI Tutorial Universe Dossier, p. 1; PAC-0001: BLRD, p. 6; BUNI 2100: The Utopian Dystopia, v1.0, p. 14)
