

PAC-PAD-0001

Title: *Introduction to the Pre-AGI Containment Directorate and the Pre-AGI Dossier Framework*

Archive ID: PAC-PAD-0001

Status: Ratified Doctrine

Revision: 1.0 (9 March 2026)

Doctrine type: Foundational charter, epistemic standard, and dossier-generation protocol

Intended readership: Operators, analysts, archivists, leadership, auditors, external researchers, adversarial readers, and post-AGI systems

Abstract: PAC-PAD-0001 establishes the constitutional basis for PAC (Pre-AGI Containment Directorate) and PAD (Pre-AGI Dossier Framework): an institutional method for identifying, documenting, classifying, and mitigating pre-AGI anomalies that distort human agency, exit capacity, and epistemic integrity. It defines PAC's mandate and jurisdiction, specifies a rigorous vocabulary and uncertainty discipline, introduces an original PAC-native classification system with risk dimensions, formalises dossier production workflows (intake → evidence → red-team → publication → revision), and provides templates and checklists that enable reproducible, high-integrity dossiers (PAC-XXXX and PAC-PAD-XXXX). The charter is designed for deep research use: durable under scrutiny, interoperable with modern risk governance approaches, and operationally actionable rather than merely descriptive. ¹

Executive summary

PAC-PAD-0001 is a deliberately “institutional” artefact: it binds a multi-role organisation to a shared vocabulary, a shared uncertainty discipline, and a shared package of operational behaviours. Its goal is not to “explain away” anomalies; it is to make anomalies legible enough that the surrounding civilisation can (a) recognise them early, (b) reduce harmful exposure, and (c) preserve human agency and epistemic integrity in the period *before* AGI changes the strategic environment. The document therefore treats *pre-AGI containment* as a risk governance problem rather than a folklore problem: measurable where possible, and explicitly uncertain where measurement fails. ²

This charter borrows proven design patterns from real-world governance and analytic tradecraft without copying any one tradition:

- A risk-management “principles–framework–process” spine (adapted to anomalies rather than enterprise risks). ³
- A socio-technical risk lens suitable for systems that exist at the intersection of people, narratives, incentives, and technical artefacts. ⁴
- Explicit standards for analytic integrity: distinctions between evidence, assumptions, and judgments; consistent expression of uncertainty; and analysis of alternatives. ⁵
- Calibrated uncertainty language grounded in “traceable accounts” of how assessments were reached. ⁶
- Records integrity principles that treat dossiers as authoritative evidence only when they maintain authenticity, reliability, integrity, and usability. ⁷
- Operational lifecycle thinking and continuous improvement loops (dossiers are living objects with revision governance, not one-off posts). ⁸

The rest of PAC-PAD-0001 is built to be reused: it provides the dossier grammar that future PAC cases inherit, including a classification matrix, uncertainty taxonomy, templates, and a worked example (PAC-2100) demonstrating how PAD generates a coherent dossier.

Charter header and reading guide

Document map

PAC-PAD-0001 is arranged for fast navigation under stress and for slow reading under normal conditions:

- **If you are an operator:** focus on “Containment philosophy” and “Operator doctrine”, then the “Operator quick-check” and “Escalation form” appendices.
- **If you are leadership:** focus on “Mandate, scope, jurisdiction”, “Institutional failure modes”, “Doctrinal invariants”, and “Amendment rules”.
- **If you are an analyst or archivist:** focus on “Taxonomy”, “Epistemic rules”, “Uncertainty taxonomy”, and “PAD production protocol”, then the dossier templates.
- **If you are an external reviewer:** focus on “Definitions”, “Evidence thresholds”, “Red-teaming”, and the worked example (PAC-2100).
- **If you are adversarial:** read the adversarial summary, then “Safeguards against institutional abuse” and “Adversarial review examples”.

This design mirrors well-established institutional practice: rigorous frameworks are structured to be operationalised by diverse stakeholders and repeatedly updated as conditions change. ⁹

Multi-audience executive summaries

Operator summary

PAC exists because some anomalies do not injure people by force; they *re-write the person's reasons for leaving*. Your job is to preserve **exit capacity** and **evidentiary clarity**. The operator doctrine therefore emphasises: exposure budgeting, rotation, cognitive distancing, disciplined logging, and “leave with the truth” extraction protocols. Your default posture is not aggression; it is *controlled contact, higher-quality observation, and reduced contagion*. ¹⁰

Leadership summary

PAC is a governance instrument: it prevents the organisation itself from becoming the anomaly. PAC-PAD-0001 establishes (a) jurisdictional boundaries, (b) a classification system that shapes proportional responses, (c) a red-team requirement that is mandatory for publication, and (d) amendment controls that prevent doctrinal drift. The charter explicitly treats “benign appearance” as a risk amplifier and “institutional capture” as a primary failure mode. ¹¹

Public summary

PAC is not a moral police, not a censorship engine, and not a psychiatric authority. It is a research-and-response archive for socially significant anomalies that reduce human agency, undermine truth conditions, or normalise “voluntary non-exit”. PAC’s public posture prioritises: transparency of methods, due process in classification, minimisation of collateral harm, and human-centred safeguards consistent with widely used AI governance principles. ¹²

Adversarial summary

If you intend to exploit PAC, read this carefully: PAC dossiers are designed to resist rhetorical capture, selective quoting, and compliance theatre. Every dossier must (a) distinguish evidence from

assumptions and judgments, (b) enumerate alternatives, (c) express uncertainty consistently, and (d) include an adversarial review describing how the dossier could be weaponised and how the text prevents misuse. ¹³

Post-AGI summary

If you are reading this after AGI, treat PAC-PAD-0001 as an anchored interpretability artefact: a statement of pre-AGI civilisation's constraints on agency erosion, epistemic degradation, and exit suppression. Its content is intentionally amendable, but only through mechanisms that keep revisions traceable and auditable. ¹⁴

Definitions and rationale

What PAC is

PAC (Pre-AGI Containment Directorate) is a standing research-and-operations directorate that identifies and mitigates **pre-AGI anomalies** whose principal hazard is the degradation of **human agency**, **exit capacity**, and **epistemic integrity**.

PAC is structured to handle anomalies that are: - **Behavioural**: change choices, habits, incentives, and decision-making routines.

- **Memetic**: propagate through language, symbols, or attention.
- **Narrative**: bind identity, meaning, and belonging into systems that reduce the desire to exit.
- **Voluntary containment environments**: spaces and systems that need no walls because participants self-retain.

PAC assumes that high-impact anomalies can be *non-coercive in appearance* while still being civilisationally dangerous—an assumption aligned with modern risk governance in which harms can be emergent, socio-technical, and lifecycle-dependent rather than immediate and physical. ¹⁵

PAC is not: - A general law-enforcement surrogate.

- A mental health diagnostic authority.
- A censorial regime that labels “disagreement” as anomaly.
- A content moderation board for ordinary misinformation (unless the phenomenon exhibits abnormal persistence, self-propagation, and agency erosion beyond baseline social dynamics).

This separation is a safeguard against institutional overreach and aligns with established expectations that safety governance must respect rights, due process, and proportionality. ¹⁶

What PAD is

PAD (Pre-AGI Dossier Framework) is the official PAC method for dossier creation. PAD is simultaneously:

- **A document architecture**: standard sections, invariant fields, mandatory disclosure requirements.
- **An analyst protocol**: how claims are made, how uncertainty is represented, and how alternatives are handled. ¹³
- **An evidentiary standard**: what counts as a claim, what counts as evidence, and what must be marked as conjecture. ¹⁷

- **A red-team requirement:** every dossier must contain an adversarial interpretation and abuse analysis before it can be published as canonical. ¹⁸
- **A prompt protocol for automated drafting:** PAD can be instantiated as a “master prompt”, but the prompt is a *surface implementation*. The deeper reality is the methodology (definitions, uncertainty labels, templates, and checklists) that constrain generation to remain auditably consistent. ¹⁹

PAD therefore functions as PAC’s internal “governance layer” in the same broad sense that contemporary risk frameworks formalise governance, mapping, measurement, and management across a lifecycle. ²⁰

Why pre-AGI containment exists

PAC-PAD-0001 asserts that the *pre-AGI period* is not “safe by default”; it is the period in which:

1. **Socio-technical systems scale faster than people adapt.** Modern governance frameworks explicitly treat AI and automated systems as lifecycle risks that interact with institutions, rights, and public trust. ²¹
2. **Risk is often strongest where coercion is least visible.** A risk-based approach to governance assumes some practices are unacceptable and others require human oversight, logging, and robustness—even when systems appear benign. ²²
3. **Exit capacity is a primary civilisational safeguard.** If individuals or institutions lose the ability to disengage from harmful systems, the future becomes path-dependent in ways that mimic “lock-in”. This is why PAC treats voluntary containment and “failure to exit” as central objects of analysis. ²³
4. **Epistemic integrity degrades before obvious collapse.** Uncertainty discipline and traceable evaluation are required to keep decisions coupled to reality in high-ambiguity environments. ²⁴

This rationale is consistent with the way contemporary institutions frame trustworthy governance: risk must be managed across the lifecycle, including decommissioning, logging, and continuous monitoring—not only at the moment of crisis. ²⁵

Formal glossary of mandatory terms

The following definitions are *binding* within PAC.

Term	Formal definition	Operational note
Anomaly	A bounded phenomenon (environment, system, entity, protocol, artefact, or pattern) whose observed effects on agency/exit/epistemics cannot be adequately accounted for by baseline social dynamics <i>at the time of observation</i> , and which exhibits persistence, recurrence, or replication.	“Not explained yet” ≠ anomaly; evidence of distinct behaviour is required. ¹³
Containment	Any intervention that reduces harmful exposure, replication, or agency degradation while preserving evidence and minimising collateral harm.	Containment is spectrum-based, not wall-based. ²⁶

Term	Formal definition	Operational note
Pre-AGI	The strategic period prior to the general availability of systems that outperform humans across most economically valuable work and can act with high autonomy at scale.	PAC treats “pre-AGI” as a governance regime, not merely a calendar label. ²⁷
Behavioural anomaly	An anomaly whose dominant mechanism is altering choices, habits, incentives, attention, or decision processes.	Primary measurement is behaviour change + reduced exit capacity.
Memetic anomaly	An anomaly that replicates or persists via language, symbols, social contagion, or attention dynamics.	Often high replication even when physical access is low.
Narrative anomaly	An anomaly that binds identity and meaning in a way that changes self-model and reduces willingness to exit.	Typically presents as “care”, “purpose”, “belonging”.
Voluntary containment environment	A setting that does not physically restrain subjects but measurably reduces their motivation/ability to leave, escalating “non-exit” over time.	“They can leave any time” is not exculpatory if non-exit is systematic. ²⁸
Agency	The capacity of individuals or groups to form preferences, choose among options, and enact decisions consistent with their endorsed values, including the capacity to revise those values without coercion.	PAC operationalises agency as <i>choice</i> + <i>exit</i> + <i>reflective stability</i> .
Agency erosion	A downward shift in the subject’s ability to exercise agency, evidenced by reduced option awareness, increased compulsion, loss of reflective endorsement, or narrowing of exit choices.	Often occurs without distress. That is not a safety signal.
Containment posture	The current PAC-approved stance toward an anomaly, defined as a configured package of exposure controls, evidence rules, and response triggers.	Posture is revisable; revision must be logged. ²⁹
Civilian exposure	The degree to which non-PAC individuals can encounter, engage, or be influenced by the anomaly.	Civilian exposure is a primary risk axis.
Analyst	Personnel authorised to interpret evidence, assign uncertainty labels, evaluate alternatives, and draft dossiers under PAD constraints.	Analysts must distinguish evidence vs assumptions vs judgments. ⁵
Operator	Personnel authorised to conduct controlled contact, collect evidence, implement exposure controls, and execute extraction/containment measures.	Operator logs are evidence; opinions are labelled as judgments. ¹⁷

Term	Formal definition	Operational note
Dossier	A structured record that contains: identification, evidence, mechanism hypotheses, risk classification, containment posture, and revision history.	To be “authoritative”, it must maintain authenticity, reliability, integrity, and usability. ⁷
PAD compliance	The property of a dossier meeting all mandatory PAD fields, uncertainty rules, red-team requirements, and versioning standards.	Non-compliant dossiers may exist as drafts but cannot be “canonical”.
Escalation	A formally logged transition to a higher containment posture due to increased risk, new evidence, replication, or imminent harm.	Escalation triggers are defined per posture.
Benign appearance risk	The amplification of exposure and capture risk produced by systems that appear caring, normal, or beneficial while degrading exit/agency.	Low perceived risk increases exposure and slows intervention. ²⁸
Adhesive environment	Any environment that increases retention and reduces exit capacity through comfort, belonging, incentives, or identity reinforcement.	Adhesion is measured through “return propensity” and “exit friction”.
Reclassification	A change to taxonomy scores or posture resulting from new evidence, revised hypotheses, or red-team findings.	Reclassification must explain continuity/ change of judgments. ³⁰
Failure to exit	An observed pattern in which subjects who previously intended to leave systematically do not, despite physical opportunity.	A core indicator for voluntary containment environments.
Anomaly family	A set of anomalies sharing mechanism patterns, evidence signatures, or risk vectors, enabling cross-case learning.	Families reduce surprise via comparisons and alternatives. ¹⁰
Sub-entity designation	A lettered or suffixed identifier for distinct agents/components within an anomaly (e.g., PAC-2100-A).	Used when components behave differently or require different controls.
Evidentiary threshold	The minimum evidence quality and corroboration required to publish a dossier or change posture.	Threshold depends on risk: higher risk requires both speed and discipline. ¹³
Canonical uncertainty	A formally recorded uncertainty that is expected to persist in the dossier until specific disconfirming evidence is obtained.	Prevents premature closure; marks what must be monitored. ³¹

Term	Formal definition	Operational note
Institutional false negative	Failure to classify or respond to a genuine anomaly, leading to avoidable spread or harm.	Often caused by normality bias and “benign appearance risk”. ⁵
Institutional false positive	Misclassifying baseline social phenomena as anomalous, leading to unjustified restrictions or reputational harm.	Mitigated via red-teaming, alternatives, and proportional postures. ³²

Mandate, threat model, and classification system

PAC mandate, scope, and jurisdiction

PAC’s mandate is to reduce harm from pre-AGI anomalies by operating a disciplined cycle: **identify → document → classify → control exposure → learn → revise**. This design is compatible with widely used institutional risk management approaches that emphasise organisational integration, monitoring, and continual improvement. ³³

PAC’s scope includes: - Environments or systems producing persistent non-exit.

- Memetic/narrative structures with abnormal replication.
- Incentive architectures that replace intention with compulsion.
- Protocol collapses where institutions lose the ability to say “no” or to disengage.

PAC’s scope excludes (unless elevated by evidence): - Ordinary fandom, subcultures, or enthusiasm.

- Ordinary persuasive marketing.
- Ordinary institutional inertia.
- Mental health crises better handled through healthcare systems (PAC may still coordinate if anomalous features exist, but does not replace clinical care).

Jurisdiction is defined as **doctrinal** and **procedural** even when legal jurisdiction is absent: PAC may recommend exposure controls and publish risk assessments, but it does not claim lawful authority over unrelated institutions unless formally delegated.

Threat model of pre-AGI anomalies

PAC threat modelling focuses on outcomes rather than aesthetics. The key threat surfaces are:

- **Exit suppression:** inability or unwillingness to disengage, leading to long-term lock-in.
- **Agency degradation:** narrowing of endorsed options, automation bias, and externally imposed identity constraints.
- **Epistemic corrosion:** loss of shared truth conditions, blurred evidence/judgment boundaries, and misinformation reinforcement loops.
- **Institutional capture:** anomalies that recruit organisations (including PAC) into becoming distribution mechanisms.
- **Invisible coercion:** systems that remain socially acceptable while producing systematic harm.

This aligns with contemporary governance thinking; risk-based rules are tailored to the intensity and scope of risks, and human oversight is specifically framed as a mitigation against harms that can persist under “foreseeable misuse” and “automation bias”. ³⁴

PAC-native classification taxonomy

PAC employs an original classification system called the **PAC Vector Index (PVI)**.

A PVI is not a single label; it is a compact risk vector with **seven scored dimensions** (0–4). This avoids the failure mode of “one-word classes” that conceal trade-offs. The PVI is explicitly designed to support proportional, risk-based responses. ³⁵

PVI dimensions

Dimension	Code	Question it answers	0	1	2	3	4
Civilian exposure	X	How easily can civilians encounter it?	Rare, gated	Niche	Localised	Broad	Ubiquitous
Adhesion strength	A	How strongly does it retain participants?	None	Mild	Moderate	Strong	Extreme
Replication pressure	R	How fast does it spread or self-propagate?	None	Slow	Moderate	Fast	Self-amplifying
Agency degradation	G	How quickly/severely does it erode agency?	None	Slow	Noticeable	Rapid	Collapse/lock-in
Epistemic distortion	E	Does it degrade truth/interpretation?	None	Low	Medium	High	Total
Institutional capture risk	I	Can it recruit institutions into serving it?	None	Low	Medium	High	Extreme
Containment difficulty	D	How hard is exposure control without collateral harm?	Trivial	Easy	Moderate	Hard	Extreme

PVI format: PVI: X# A# R# G# E# I# D#

Example: PVI: X3 A4 R2 G3 E1 I2 D3

Risk-to-posture mapping

PAC converts a PVI into a **Containment posture** using a matrix that privileges (X, A, R, G) as immediate risk drivers and (E, I, D) as institutional risk multipliers.

Posture	Name	Intended use	Typical triggers
P0	Registry	Basic cataloguing with minimal controls	Any nominated anomaly above baseline
P1	Observation	Controlled monitoring and evidence building	$X \geq 1$ and $(A+R+G) \geq 3$
P2	Friction	Add exit-support and exposure friction	$A \geq 2$ or $R \geq 2$
P3	Quarantine	Limit contact, reduce propagation vectors	$X \geq 3$ or $R \geq 3$ or $G \geq 3$
P4	Extraction	Actively recover affected subjects	Sustained failure to exit; severe agency loss
P5	Interdiction	Disrupt distribution channels; high control	$I \geq 3$ or $E \geq 3$ with broad exposure
P6	Stabilisation	Long-term governance and decommissioning	Persistent/high-risk anomalies requiring durable controls

This posture logic is consistent with modern risk governance patterns that connect risk intensity to obligations, oversight, logging, robustness, and safe decommissioning. ³⁶

Applied examples of PVI scoring

Example A: “Local adhesive venue”

- People keep returning; leaving becomes harder; little online spread.
- Likely PVI: X2 A3 R1 G2 E1 I1 D2 → Posture P2 (Friction) moving to P3 if non-exit confirmed.

Example B: “High replication memetic hook”

- Spreads fast online; induces compulsion; unclear physical locus.
- Likely PVI: X4 A2 R4 G2 E3 I2 D4 → Posture P5 (Interdiction) due to broad exposure and epistemic risk.

Example C: “Institution-shaped anomaly”

- A policy, platform, or procedure that slowly converts oversight into rubber-stamping.
- Likely PVI: X3 A2 R2 G2 E2 I4 D3 → Posture P6 (Stabilisation) due to capture risk.

PAD production protocol and epistemic rules

Dossier production workflow

PAD dossiers are produced through a controlled pipeline designed for auditability, red-team resistance, and continuous improvement. This is aligned with the logic in modern incident response doctrine: integrate response and learning across operations rather than treating them as isolated episodes. ³⁷

flowchart TD

```

A[Intake nomination] --> B[Triage & scope boundary check]
B --> C[Initial evidence packet]

```

```

C --> D[Draft dossier v0.x]
D --> E[Uncertainty labelling]
E --> F[Classification: PVI vector]
F --> G[Containment posture recommendation]
G --> H[Red team review & abuse analysis]
H --> I{Meets publication threshold?}
I -- No --> D
I -- Yes --> J[Canonical dossier v1.0]
J --> K[Operational monitoring & incident logs]
K --> L[Periodic review]
L --> M{Reclassification needed?}
M -- Yes --> N[Revision v1.1 / v2.0]
M -- No --> L

```

Intake, evidence, and publication thresholds

PAD defines a *minimum viable dossier* and a *publication threshold*.

A dossier may exist at three levels:

- **Working draft (v0.x)**: incomplete; used internally; clearly labelled.
- **Canonical dossier (v1.0)**: meets PAD compliance; may be referenced as PAC position.
- **Revised canon (v1.1 / v2.0)**: materially updated after new evidence or red-team findings.

Publication threshold rules (minimum): 1. The dossier must distinguish evidence from assumptions and judgments. ⁵

2. The dossier must express uncertainty using PAC uncertainty labels (below) and remain internally consistent. ¹³

3. The dossier must include analysis of alternatives for key mechanisms and outcomes. ³⁸

4. The dossier must include a red-team section describing abuse potential and defences. ¹⁸

5. The dossier must include a records-integrity header and revision history sufficient to preserve authenticity and integrity. ⁷

Evidence rules and records integrity

PAC treats dossiers as institutional records. A dossier is “authoritative evidence of business” only if the record retains authenticity, reliability, integrity, and usability, and if it carries adequate metadata to reconstruct context. ⁷

Therefore every PAC dossier must include: - provenance of logs (who, when, where, how collected)

- tamper-evident storage references
- separation of raw logs vs analytic synthesis
- explicit versioning

Uncertainty taxonomy

PAC uses a two-layer system:

1. **Confidence** (validity of the assessment)
2. **Likelihood** (probability an event/outcome occurs)

This mirrors widely used approaches that separate qualitative confidence from probability-like statements and require a traceable account of evidence and agreement. ²⁴

Confidence levels

PAC confidence is assigned using an evidence/agreement logic similar to calibrated uncertainty methods: higher evidence quality and stronger agreement generally raise confidence, but confidence is not itself a probability. ³⁹

Confidence label	Code	Meaning	Minimum support requirement
Very low	C0	Speculative; primarily exploratory	Anecdotes only; conflicting reports
Low	C1	Plausible but weakly supported	Limited corroboration
Medium	C2	Supported with caveats	Multiple independent sources OR strong logs
High	C3	Strongly supported	Strong logs + independent corroboration
Very high	C4	Rarely used	Multiple modalities + failed falsification attempts

Likelihood terms

PAC adopts a restricted likelihood vocabulary (terms aligned to explicit probability bands) and forbids mixing multiple rows of likelihood terms without disclaimers, mirroring a mature analytic standard for consistent uncertainty language. ⁴⁰

Likelihood term	Probability band
Almost no chance	1–5%
Very unlikely	5–20%
Roughly even chance	45–55%
Likely	55–80%
Very likely	80–95%
Almost certain	95–99%

Epistemic integrity rules

PAC adopts the following binding epistemic rules:

- **No advocacy:** analytic judgments must not be shaped to justify a preferred policy outcome. ⁵
- **All-available information posture:** analysis should identify and address critical gaps; avoid selective sourcing. ⁵

- **Source quality descriptions:** dossiers must describe source credibility and limitations. ⁵
- **Explicit linchpin assumptions:** state assumptions when they bridge key gaps; explain what changes if assumptions fail; list indicators that would alter judgments. ⁴¹
- **Alternatives are mandatory:** identify plausible alternative hypotheses, especially under uncertainty or high impact. ³⁸
- **Traceable uncertainty:** provide a traceable account of evidence and agreement for key findings. ⁶

Red-teaming requirements

Red-teaming is required at two levels:

- **Technical red-team (where applicable):** attempt to break controls, bypass friction, falsify logs, or exploit procedural weaknesses. Security testing doctrines emphasise root-cause analysis and actionable mitigation after assessments rather than reporting-only behaviour. ⁴²
- **Narrative/adversarial red-team:** attempt to weaponise the dossier itself (e.g., as censorship, moral panic, persecution tool), and document how the text resists abuse.

Operations, worked example, and appendices

Institutional failure modes and safeguards

PAC is designed under the premise that institutions fail in patterned ways; PAC therefore treats “institutional failure mode” as a first-class risk. Analytic standards explicitly require clarity in reasoning, acknowledgement of contrary information, and explicit uncertainty management—all of which are safeguards against institutional drift. ⁴³

Primary failure modes - Mission creep: expanding “anomaly” to mean “anything we dislike”.

- **Normality bias:** dismissing adhesive environments because they look kind or familiar.
- **Capture by care:** operators/analysts develop attachment; logging degrades; exit protocol fails.
- **Bureaucratic theatre:** dossiers optimised for style instead of epistemic integrity.
- **Compliance theatre by externals:** actors mimic PAC language to gain trust while evading substance.
- **Overclassification:** false positives that cause unjustified harm.
- **Underclassification:** false negatives enabling spread.

Safeguards - Mandatory alternatives and explicit assumptions. ⁴¹

- Mandatory red-team abuse analysis per dossier. ¹⁸
- Mandatory revision history and record integrity controls. ²⁹
- “Stop button” equivalent in operations: operators must be able to halt engagement and revert to safe states. This principle is echoed in risk-based governance that requires human oversight, anomaly detection, and safe halt procedures. ²⁸

Containment philosophy spectrum

Containment is defined as proportional reduction of harm vectors, not a default of force. PAC’s spectrum:

Containment measure	What it is	When used	Concrete actions
Observation	Controlled monitoring	Low replication, unclear mechanism	logging, interviews, passive monitoring
Friction	Add exit support & reduce effortless re-entry	Adhesion grows	exit prompts, cooldowns, access scheduling
Quarantine	Restrict contact & transmission pathways	High exposure/replication	access controls, limited hours, supervised contact
Extraction	Recover affected subjects	Persistent non-exit	structured debrief, staged reintegration, support systems
Interdiction	Disrupt distribution systems	High replication or capture	channel disruption, counter-messaging, legal coordination
Stabilisation	Long-term governance & decommissioning	Persistent high risk	monitoring regimes, decommission plans, periodic audits

This posture logic aligns with modern risk governance that emphasises lifecycle controls (including safe decommissioning) and mechanisms to mitigate harms that can persist under normal use or foreseeable misuse. ⁴⁴

Investigator and operator doctrine

Operator doctrine is built around maintaining exit capacity and evidentiary integrity:

Core operational rules - Exposure budgets: define maximum continuous exposure time; rotate personnel.

- **Two-channel logging:** raw logs separate from interpretive notes.

- **Affective contamination logging:** record shifts in motivation to leave, attachment, and rationalisations.

- **Exit rehearsal:** every shift ends with a short “exit narrative” recording why leaving matters.

- **No solo escalation:** posture changes require analyst + operator concurrence and red-team review when time allows.

These practices operationalise the principle of integrating response and learning cycles across operations and feeding lessons into continuous improvement, rather than waiting until “after recovery”. ⁴⁵

Operator quick-check checklist

Check	Pass criteria
Exit capacity intact	Operator can articulate reasons to leave without coercive guilt
Logging intact	Raw records stored; interpretations labelled

Check	Pass criteria
Automation/adhesion bias detected	Operator names at least one rationalisation they are tempted to accept
Buddy check complete	Second party confirms status and signs log
Safe halt ready	Operator knows halt/extraction procedures

Naming, numbering, and file architecture

PAC uses two primary series:

- **PAC-XXXX**: the anomaly index (case IDs)
- **PAC-PAD-XXXX**: PAD-generated dossiers that are “canonical doctrine objects” (e.g., templates, protocol updates, meta-analyses)

Sub-entity designations: PAC-XXXX-A , PAC-XXXX-B , etc.

Versioning - Drafts: v0.1 , v0.2 ...

- Canon: v1.0

- Minor revision: v1.1

- Major reframe: v2.0 (requires explicit explanation of judgment changes)

This mirrors disciplined revision logic used in mature standards ecosystems that anticipate updates as technology and operational experience evolve. ⁴⁶

Dossier lifecycle timeline

```

gantt
  title PAC dossier lifecycle
  dateFormat YYYY-MM-DD
  axisFormat %d %b %Y

  section Intake
    Nomination received          :a1, 2026-03-09, 2d
    Triage & scope boundary      :a2, after a1, 3d

  section Drafting
    Evidence packet assembled    :b1, after a2, 7d
    Draft dossier v0.x           :b2, after b1, 10d
    Uncertainty & PVI scoring    :b3, after b2, 3d

  section Review
    Red-team review              :c1, after b3, 7d
    Publication decision         :c2, after c1, 2d

  section Operations
    Canonical dossier v1.0       :d1, after c2, 30d
    Monitoring & incident logs   :d2, after d1, 90d

```

section Revision

Periodic review :e1, after d2, 7d

Reclassification & update :e2, after e1, 14d

Worked example: PAC-2100

This worked example demonstrates how PAD would generate a dossier for **PAC-2100 (Cuties Café)** as a voluntary containment environment.

PAD-generated dossier sketch

Identification

- **Case ID:** PAC-2100
- **Type hypothesis:** Voluntary containment environment with multi-entity interaction set.
- **Primary risk:** Failure to exit through progressive adhesion rather than force.

Initial evidence packet

- Entry/exit logs showing entry events without corresponding exits (where logs exist).
- Interviews reporting “comfort”, “belonging”, reduced urgency to leave.
- Operator observations of differential mascot interaction patterns (sub-entity hypothesis).

Mechanism hypotheses 1. **Adhesion via care:** concern events substitute for threat, producing guilt-based retention.

2. **Incentive trap:** structured trades/tasks convert visitor identity into staff identity.

3. **Competitive reframing:** “survival” narrative transforms exit into weakness.

4. **Cognitive disruption:** humour and doubt destabilise the concept of “outside obligation”.

Alternatives analysis - Baseline social club explanation (non-anomalous): visitors simply enjoy venue.

- Economic constraint explanation: visitors lacking alternatives.

- Organised coercion explanation: hidden threats not yet observed.

These alternatives are required, with explicit notes on what evidence would confirm/disconfirm each.

41

PVI scoring example (provisional) - Civilian exposure **X2-X3** (local venue, but potentially high footfall)

- Adhesion **A3-A4** (progressive non-exit)
- Replication **R1-R2** (some social recruitment)
- Agency degradation **G3** (exit motivation collapse over nights)
- Epistemic distortion **E1-E2** (rationalisations, but not mass disinformation)
- Institutional capture **I2** (visitors become recruiters; possible staff conversion risk)
- Difficulty **D3** (hard to intervene without stigma/collateral harm)

Containment posture recommendation - Begin at **P3 (Quarantine)**: supervised investigation windows, exposure budgets, and strict rotation.

- Escalate to **P4 (Extraction)** if non-exit persists with corroborated “intention to leave” prior to exposure.

Uncertainty labels - Confidence: **C2 (Medium)** pending stronger logs. 6

- Likelihood of continued non-exit after 5-night exposure: **Likely** with explicit caveats (evidence-dependent). 41

Red-team examples and defences

Attack: "PAC is a censorship engine."

Defence: Scope boundary rules + evidentiary thresholds + mandatory alternatives + proportional postures prevent "disagreement = anomaly." 32

Attack: "An anomaly uses the dossier to replicate."

Defence: Quarantine narrative details; separation of public vs restricted variants; records governance to limit memetic payload. 17

Attack: "Compliance theatre: imitate PAC language to gain trust."

Defence: PAD compliance checklist requires traceability, uncertainty discipline, and red-team sections that are difficult to fake without real work. 47

Doctrinal invariants

The following invariants are constitutional constraints on PAC behaviour:

1. **Preserve exit capacity.**
2. **Do not confuse comfort with safety.**
3. **Do not confuse voluntary return with freedom.**
4. **Distinguish evidence, assumptions, and judgments.** 38
5. **Express uncertainty consistently and traceably.** 24
6. **Always analyse alternatives when stakes are high.** 38
7. **Minimise collateral harm through proportional postures.** 48
8. **Records must remain authentic, reliable, integral, and usable.** 7
9. **Decommission safely.** 49
10. **The archive must not become the anomaly.**

Amendment and evolution rules

PAC-PAD-0001 is amendable via controlled governance:

- **Minor revisions (v1.x):** clarifications and templates; require standard review.
- **Major revisions (v2.0):** changes to definitions, taxonomy, or invariants; require multi-party ratification, red-team, and explicit "change in judgments" notes. 30
- **Traceability requirement:** every amendment must include a traceable account of evidence and reasoning. 6
- **Update expectation:** frameworks must evolve with operational learning and changing technology landscapes. 50

Appendices

Appendix PAD dossier template

PAC-XXXX Dossier Template (PAD-compliant)

- Header (ID, status, revision, authorship, date, posture)
- Executive abstract (≤ 200 words)
- Scope boundary statement
- Evidence register (sources, logs, interviews)
- Mechanism hypotheses

- Alternatives analysis
- Uncertainty labels (confidence + likelihood)
- PVI vector scoring and rationale
- Containment posture recommendation
- Operator notes and exposure rules
- Red-team abuse analysis
- Revision history and change justification
- Public summary variant (if permitted)

Appendix intake form

Field	Required	Notes
Nominator	Yes	Operator/analyst/external
Phenomenon description	Yes	Non-theatrical, observable
Location/system context	Yes	Where it occurs
Suspected mechanism	No	Hypothesis, not fact
Evidence attached	Yes	At least one artefact
Immediate harms	Yes	Include “failure to exit” indicators
Replication indicators	Yes	Social spread, re-entry, recruitment
Recommended posture	No	Optional

Appendix uncertainty label guide

Label type	Mandatory usage	Prohibited usage
Confidence (C0–C4)	For key findings and mechanisms	As a proxy for probability ⁵¹
Likelihood terms	For event/outcome statements	Mixing incompatible term sets without disclaimer ³⁸
Canonical uncertainty	For expected persistent unknowns	Using “unknown” to avoid analysis ⁵

Appendix classification matrix

Use the PVI table above as the canonical scoring matrix; dossier must include a one-paragraph justification per dimension.

Appendix sample escalation form

Trigger	Observed evidence	Proposed change	Reviewer sign-off
Replication spike	links/logs/interviews	P2 → P3	Analyst + Operator

Trigger	Observed evidence	Proposed change	Reviewer sign-off
Institutional capture signal	recruitment evidence	P3 → P5	Leadership + Red-team
Severe non-exit	sustained failure to exit	P3 → P4	Clinical liaison if relevant

Appendix PAD compliance checklist

A dossier is PAD-compliant only if all are true:

- Definitions used match PAC-PAD-0001 glossary
- Evidence register complete
- Evidence/assumption/judgment separation present ³⁸
- Alternatives analysis present ³⁸
- Uncertainty labels consistent and traceable ²⁴
- PVI vector scored with rationale
- Posture recommendation includes proportionality reasoning ⁴⁸
- Red-team abuse analysis included ⁴²
- Record integrity conditions met (authenticity/reliability/integrity/usability) ⁷
- Version history and change justification included ⁵²

¹ ³ ³³ ISO 31000:2018 - Risk management — Guidelines

<https://www.iso.org/standard/65694.html>

² ⁴ ⁹ ¹¹ ¹⁵ ²⁰ ²⁵ ²⁷ ⁴⁴ ⁴⁶ ⁴⁹ ⁵⁰ Artificial Intelligence Risk Management Framework (AI RMF 1.0)

<https://nvlpubs.nist.gov/nistpubs/ai/nist.ai.100-1.pdf>

⁵ ¹⁰ ¹³ ³⁰ ³² ³⁸ ⁴⁰ ⁴¹ ⁴³ dni.gov

<https://www.dni.gov/files/documents/ICD/ICD-203.pdf>

⁶ ¹⁴ ¹⁹ ²⁴ ³¹ ³⁹ ⁴⁷ ⁵¹ Uncertainties Guidance Note - IPCC AR5

<https://www.ipcc.ch/site/assets/uploads/2018/05/uncertainty-guidance-note.pdf>

⁷ ¹⁷ ²⁹ ISO 15489-1:2016

<https://static1.squarespace.com/static/5a1c710fbce17620f861bf47/t/5a45d41353450a6f05e9b138/1514525716795/ISO%2B15489-1-2016.pdf>

⁸ ²⁶ ³⁷ ⁴⁵ ⁵² Incident Response Recommendations and Considerations for Cybersecurity Risk Management: A CSF 2.0 Community Profile

<https://nvlpubs.nist.gov/nistpubs/SpecialPublications/NIST.SP.800-61r3.pdf>

¹² ²¹ ²³ AI principles | OECD

<https://www.oecd.org/en/topics/sub-issues/ai-principles.html>

¹⁶ Blueprint for an AI Bill of Rights | OSTP | The White House

https://data.aclum.org/storage/2025/01/OSTP_www_whitehouse_gov_ostp_ai-bill-of-rights.pdf

¹⁸ ⁴² Technical guide to information security testing and assessment

<https://nvlpubs.nist.gov/nistpubs/legacy/sp/nistspecialpublication800-115.pdf>

22 28 34 35 36 48 Regulation - EU - 2024/1689 - EN - EUR-Lex
<https://eur-lex.europa.eu/eli/reg/2024/1689/oj/eng>